

DexVision / Hand2Bot

Vision-based dexterous teleoperation, learning feasibility, comprehensive skill data, and future language-guided orchestration

CURRENT POSITION: Levels 1–2 complete. Level 3 is active and asks whether learning works on the existing data. Level 4 is the major dataset and skill-library campaign.

This overview is generated from docs/project_overview.md. CURRENT_STATUS.md remains the checkpoint source of truth.

Project identity

DexVision is a vision-based dexterous teleoperation, robot-learning dataset, and reusable manipulation-skill project. A camera tracks a human hand, the motion is retargeted to a Shadow Hand in MuJoCo, demonstrations are recorded and validated, and PyTorch policies learn bounded low-level skills.


The eventual language model is a planner, not the robot controller. It should turn requests into typed calls to already-qualified skills. A deterministic supervisor validates parameters and world state, enforces safety and timeouts, and reports structured success or failure.

Current status

Levels 1 and 2 are complete. The repository has live full-hand/base/wrist teleoperation, resettable manipulation tasks, demonstration recording and semantic replay, quality filtering and relabeling, dataset summaries, three retargeting methods, a full retargeting benchmark, and an immutable Level 2 dataset release tracked with Git LFS.

Level 3 is active. Its purpose is to test whether the existing narrow datasets support useful behavior-cloned policies. The next checkpoint is Level 3.1, the goal-conditioned per-skill dataset loader.

Architecture



The continuous skill policy receives state plus a typed goal and outputs bounded base/wrist/finger actions. Visual perception produces stable object ids and poses. A language model may help select or disambiguate objects, but it does not emit high-rate actuator commands.

Roadmap

Level 1 — Teleoperation (complete)

OpenCV and MediaPipe track a human hand. Calibrated image motion, hand scale, palm rotation, and finger features control the MuJoCo Shadow Hand through the full Level 1.13 action space.

Level 2 — Demonstrations and data infrastructure (complete)

The system records, replays, validates, relabels, filters, summarizes, and benchmarks demonstrations. The released manipulation data includes 55 `reach_touch_target`, 55 `button_press`, and 101 `push_cube_to_target` clean successes. The archive is immutable and accompanied by a checksum and manifest.

Level 3 — Learning feasibility (active)

Level 3 builds a deterministic PyTorch loader, a small goal-conditioned MLP, a reproducible behavior-cloning loop, and closed-loop MuJoCo evaluation. It then checks whether the same approach works for reach, button press, and cube push, diagnoses the role of data quality and action fields, and adds a temporal baseline only if measured failures justify it.

Passing Level 3 proves the learning pipeline on the Level 2 distributions. It does not establish cross-session, cross-object, visual, or open-world generalization.

Level 4 — Comprehensive dataset and qualified skills (planned)

Level 4 is the major data haul and skill-library phase. New data must carry genuine session provenance, broader object and goal variation, synchronized visual labels when used, explicit failures, and corrective/recovery episodes. The planning floor is 100 accepted episodes per core skill family across at least three genuine sessions, subject to a documented pilot that can revise the count.

Core skill families are:

```

reach/approach
grasp/hold/lift
transport/place/release
push/slide
press/rotate
regrasp/drop recovery

```

The successful Level 3 approach is retrained and evaluated on session-held-out and condition-held-out data. Only policies that meet frozen closed-loop gates enter the skill registry and supervised executor.

Visual grounding starts with simulator truth and rendered boxes, masks, and poses, then a conventional detector/tracker. A compact local VLM may be tested for semantic object selection; metric localization and safety remain separate.

Level 5 — Robustness and portfolio polish (planned)

The previous Level 4 roadmap is preserved here: README and architecture polish, results tables, demo materials, robust CLI/config handling, reproducibility and CI, troubleshooting, optional dataset export, and carefully labeled advanced extensions.

Level 6 — Language-guided orchestration (future)

An LLM or deterministic planner consumes symbolic world state and versioned skill cards. It creates typed plans; a deterministic supervisor validates and executes them. Scripted plans and mock skills validate the contract before learned policies are introduced one at a time.

Diverse pilot tasks

The final evaluation is not a single sandwich demo. Level 4 must validate at least three materially different scripted pilots through the same typed skill interface:

1. Sort and pack varied objects into matching bins or receptacles.
2. Clear a workspace into a tray and recover from at least one miss.
3. Press selected buttons and rotate a dial in a requested sequence.
4. Assemble a simple sandwich from rigid ingredient proxies on a plate.
5. Set a place with a cup, plate, and utensil proxy.

Rigid-proxy sandwich assembly is feasible because it mainly composes grasp, transport, place, and release. Real tomato cutting is not a core target for the current hand-only setup: deformable/fracture physics, blade contact, force control, tool safety, and broader arm motion make it a different research problem. A guided pre-scored or pre-segmented rigid cutting proxy may be an explicitly labeled later experiment.

Dataset release policy

Editable operator data stays under ignored [data/demos/](#). Immutable bounded releases use Git LFS with a manifest and SHA-256 checksum. The Level 2 archive must never be overwritten. If Level 4 images exceed repository hosting quotas, Git should still track manifests, splits, checksums, and retrieval instructions while a versioned external artifact store holds the immutable payload.

Honest project claim

DexVision currently demonstrates a complete teleoperation and dataset engine. Level 3 will determine whether its existing data can support autonomous policies. Level 4 will determine whether those methods can become a broad, qualified skill library. Language-guided multi-skill behavior remains future work until those evidence gates pass.